

External Reading Collection - Personal Knowledge Base

External Reading: Retrieval-Augmented Generation - Page 1

Retrieval-Augmented Generation (RAG) combines a retrieval module with a text generation model. The retrieval module searches an external knowledge base for passages relevant to the user query.

Key benefit: answers are grounded in source documents, which reduces hallucination and enables citation of sources.

External Reading Collection - Personal Knowledge Base

External Reading: Retrieval-Augmented Generation - Page 2

A vector database stores text embeddings and supports fast similarity search. Popular choices include Qdrant, Milvus and Weaviate. Qdrant supports payload filtering, which allows the system to restrict search results by metadata fields such as domain, category, topic and status.

Chunking strategy matters: chunks that are too small lose context, while chunks that are too large dilute relevance.

External Reading Collection - Personal Knowledge Base

External Reading: Retrieval-Augmented Generation - Page 3

Recommended default parameters for a personal knowledge base:
chunk size 500 to 800 characters, overlap 80 to 150 characters,
and top-k between 3 and 5. Always evaluate retrieval quality
with a question set before tuning any further.

Source: external articles collected in August 2026.